

```
In [1]: # Import necessary libraries for data analysis and visualization
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import plotly.graph_objects as go
```

```
In [2]: # Read the Oktoberfest dataset from a CSV file into a Pandas DataFrame

df = pd.read_csv('oktoberfestgesamt19852022.csv')
```

```
In [3]: # Display the first few rows of the DataFrame to get an initial view of the data
# This helps in understanding the structure and contents of the dataset

df.head()
```

```
Out[3]:
```

	jahr	dauer	besucher_gesamt	besucher_tag	bier_preis	bier_konsum	hendl_preis	hendl_konsum
0	1985	16	7.1	444	3.20	54541	4.77	629520
1	1986	16	6.7	419	3.30	53807	3.92	698137
2	1987	16	6.5	406	3.37	51842	3.98	732859
3	1988	16	5.7	356	3.45	50951	4.19	720139
4	1989	16	6.2	388	3.60	51241	4.22	775674

```
In [4]: # Remove the 'besucher_tag' column from the DataFrame
# axis=1 indicates that we're dropping a column (not a row).
# inplace=True means the DataFrame is modified in place, and the result is saved in 'df'

df.drop('besucher_tag', axis=1, inplace=True)
```

```
In [5]: # Translate the column names to english
# Create a dictionary 'german_to_english' to map German column names to English equivalents
german_to_english = {'hendl_konsum':'chicken_comsump', # Rename 'hendl_konsum' to 'chicken_comsump'
                      'hendl_preis':'chicken_price',   # Rename 'hendl_preis' to 'chicken_price'
                      'bier_konsum':'bier_consump',     # Rename 'bier_konsum' to 'bier_consump'
                      'jahr':'year',                   # Rename 'jahr' to 'year'
                      'dauer':'duration',              # Rename 'dauer' to 'duration'
                      'besucher_gesamt':'visitors_total', # Rename 'besucher_gesamt' to 'visitors_total'
                      'bier_preis':'bier_price'}         # Rename 'bier_preis' to 'bier_price'

# Use the df.rename() method to rename columns based on the 'german_to_english' mapping
# inplace=True means the changes are applied directly to the DataFrame 'df'.

df.rename(columns=german_to_english, inplace = True)

# Display the first few rows of the DataFrame after renaming the columns

df.head()
```

```
Out[5]:
```

	year	duration	visitors_total	bier_price	bier_consump	chicken_price	chicken_comsump
0	1985	16	7.1	3.20	54541	4.77	629520
1	1986	16	6.7	3.30	53807	3.92	698137
2	1987	16	6.5	3.37	51842	3.98	732859
3	1988	16	5.7	3.45	50951	4.19	720139
4	1989	16	6.2	3.60	51241	4.22	775674

```
In [6]: # Get the shape of the DataFrame 'df'
# The 'shape' attribute returns a tuple representing the dimensions of the DataFrame.
# The first element of the tuple is the number of rows, and the second element is the nu

df.shape
```

Out[6]: (36, 7)

```
In [7]: # Display information about the DataFrame 'df'
# The 'info()' method provides a concise summary of the DataFrame's structure and conten
# It includes details such as the number of non-null entries, data types of columns, and

df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 36 entries, 0 to 35
Data columns (total 7 columns):
#   Column                Non-Null Count  Dtype
---  -
0   year                  36 non-null    int64
1   duration              36 non-null    int64
2   visitors_total        36 non-null    float64
3   bier_price            36 non-null    float64
4   bier_consump          36 non-null    int64
5   chicken_price         36 non-null    float64
6   chicken_comsump       36 non-null    int64
dtypes: float64(3), int64(4)
memory usage: 2.1 KB
```

```
In [8]: # Generate summary statistics of the DataFrame 'df'
# The 'describe()' method computes various summary statistics for numerical columns in t
# These statistics include count, mean, standard deviation, minimum, and quartiles.

df.describe()
```

Out[8]:

	year	duration	visitors_total	bier_price	bier_consump	chicken_price	chicken_comsump
<b>count</b>	36.000000	36.000000	36.000000	36.000000	36.000000	36.000000	36.000000
<b>mean</b>	2002.555556	16.305556	6.297222	7.053333	62476.611111	7.768889	564746.138889
<b>std</b>	10.635371	0.624246	0.398200	2.770215	10196.196093	2.793060	129222.529768
<b>min</b>	1985.000000	16.000000	5.500000	3.200000	48698.000000	3.920000	313636.000000
<b>25%</b>	1993.750000	16.000000	5.975000	4.845000	53510.750000	5.317500	481139.250000
<b>50%</b>	2002.500000	16.000000	6.350000	6.750000	61467.500000	8.130000	515646.000000
<b>75%</b>	2011.250000	16.000000	6.500000	9.170000	71085.000000	9.890000	685465.750000
<b>max</b>	2022.000000	18.000000	7.100000	13.450000	79225.000000	13.960000	807710.000000

```
In [9]: # Check if there are any missing values (NaN or null values) in the DataFrame 'df'

df.isnull().sum().any()
```

Out[9]: False

```
In [10]: # Get a list of column names in the DataFrame 'df'

df.columns
```

```
Index(['year', 'duration', 'visitors_total', 'bier_price', 'bier_consump',
```


```
Out[10]:      'chicken_price', 'chicken_consump'],  
          dtype='object')
```

```
In [11]: # Sort the DataFrame 'df' by the 'beer_consump' column in descending order  
         # and display the top (largest) values  
  
df['bier_consump'].sort_values(ascending=False).head()
```

```
Out[11]: 26      79225  
        33      78705  
        34      78502  
        32      77836  
        28      77031  
        Name: bier_consump, dtype: int64
```

```
In [12]: # Extract 'year' and 'beer_consump' columns from the DataFrame 'df'  
  
year = df['year']  
bier_consumption = df['bier_consump']  
  
# Create a Plotly figure for the line plot  
  
fig = go.Figure()  
  
# Add a line and markers trace for bier consumption data  
  
fig.add_trace(go.Scatter(x=year, y=bier_consumption,  
                        mode='lines+markers',  
                        name='Bier', line=dict(color='red')))  
  
# Customize the layout of the plot  
  
fig.update_layout(title='Consumption of beer from 1985 to 2022 in Oktoberfest',  
                  xaxis_title='Year',  
                  yaxis_title='Consumption(Hectoliters)')
```

## Consumption of beer from 1985 to 2022 in Oktoberfest



```
In [13]: # Extract 'year' and 'chicken_comsump' columns from the DataFrame 'df'

year = df['year']
chicken_consumption = df['chicken_comsump']

# Create a Plotly figure for the line plot

fig = go.Figure()


# Add a line and markers trace for chicken consumption data

fig.add_trace(go.Scatter(x=year, y=chicken_consumption,
                          mode='lines+markers',
                          name='Chicken',
                          line=dict(color='orange'))))

# Customize the layout of the plot

fig.update_layout(title='Consumption of Chicken from 1985 to 2022 in Oktoberfest',
                  xaxis_title='Year',
                  yaxis_title='Consumption(Units)')
```

### Consumption of Chicken from 1985 to 2022 in Oktoberfest



```
In [14]: # Create a new figure with a specified size

plt.figure(figsize=(15, 8))

# Create a line plot for bier prices using Seaborn
```

```

sns.lineplot(data=df, x='year', y='bier_price', linewidth=3, label='Bier Price')

# Create a line plot for chicken price on the same plot
sns.lineplot(data=df, x='year', y='chicken_price', linewidth=3, label='Chicken Price')

# Set the title and axis labels

plt.title('Beer and Chicken Prices Over the Years', fontsize=20)
plt.xlabel('Year', fontsize=16)
plt.ylabel('Price (Euro)', fontsize=16)

# Display gridlines

plt.grid(True)


# Display a legend with labels for the two lines

plt.legend(fontsize=16)

# Show the plot

plt.show()

```



```


In [15]: year = df['year']
chicken_price = df['chicken_price']
bier_price = df['bier_price']

fig = go.Figure()
fig.add_trace(go.Scatter(x=year, y=chicken_price,
                        mode='lines+markers',
                        name='Chicken'))
fig.add_trace(go.Scatter(x=year, y=bier_price,
                        mode='lines+markers',
                        name='Bier'))

fig.update_layout(title='Price Comparison between Chicken and Bier prices Over the years',
                  xaxis_title='Year',
                  yaxis_title='Price (Euros)')

```

## Price Comparison between Chicken and Bier prices Over the years



```
In [16]: # Create a new column 'total_visitors_million' in the DataFrame 'df'
# This column represents the total number of visitors in millions

# Multiply the 'visitors_total' column by 1,000,000 to convert it to millions

df['total_visitors_million'] = df['visitors_total']*1_000_000

# Display the first few rows of the DataFrame after adding the new column

df.head()
```

```
Out[16]:
```

	year	duration	visitors_total	bier_price	bier_consump	chicken_price	chicken_comsump	total_visitors_million
0	1985	16	7.1	3.20	54541	4.77	629520	7100000.0
1	1986	16	6.7	3.30	53807	3.92	698137	6700000.0
2	1987	16	6.5	3.37	51842	3.98	732859	6500000.0
3	1988	16	5.7	3.45	50951	4.19	720139	5700000.0
4	1989	16	6.2	3.60	51241	4.22	775674	6200000.0

```
In [17]: # Calculate the liters of beer per person by dividing 'bier_consump' by 'total_visitors_
year = df['year']
beer_per_person = df['bier_consump']*100 / df['total_visitors_million']

# Create a Plotly figure for the line plot

fig = go.Figure()
```


```
# Add a line and markers trace for beer consumption per person data

fig.add_trace(go.Scatter(x=year, y=beer_per_person,
                        mode='lines+markers',
                        name='bier/person'))

# Customize the layout of the plot

fig.update_layout(title='Bier Consumption per person over the years',
                  xaxis_title='Year',
                  yaxis_title='Liters')
```

## Bier Consumption per person over the years



```
In [18]: # Calculate the correlation between beer consumption and beer price

beer_correlation = df['bier_consump'].corr(df['bier_price'])

# Calculate the correlation between chicken consumption and chicken price

chicken_correlation = df['chicken_comsump'].corr(df['chicken_price'])

# Print the calculated correlations

print(f"Correlation between Beer Consumption and Beer Price: {beer_correlation:.2f}")
print(f"Correlation between Chicken Consumption and Chicken Price: {chicken_correlation:.2f}")

Correlation between Beer Consumption and Beer Price: 0.89
Correlation between Chicken Consumption and Chicken Price: -0.86
```

```
In [19]: # Create subplots for beer and chicken correlations
# 1 row, 2 columns of subplots, figsize is set to (12, 4)

fig, (ax1, ax2) = plt.subplots(1, 2, figsize=(12, 4))
```

```

# Create a scatter plot for beer consumption vs. beer price on the first subplot (ax1)

ax1.scatter(df['bier_consump'], df['bier_price'], alpha=0.5)
ax1.set_xlabel('Bier Consumption')
ax1.set_ylabel('Bier Price')
ax1.set_title('Bier Consumption vs. Bier Price')

# Create a scatter plot for chicken consumption vs. chicken price on the second subplot

ax2.scatter(df['chicken_comsump'], df['chicken_price'], alpha=0.5)
ax2.set_xlabel('Chicken Consumption')
ax2.set_ylabel('Chicken Price')
ax2.set_title('Chicken Consumption vs. Chicken Price')

# Adjust subplot layout to prevent overlap

plt.tight_layout()

# Show the subplots

plt.show()

```

